

# CSCI 596: HW 6

M. Oneeb H. Khan

October 25, 2021

## 1 Pair-Distribution Computation with CUDA

### 1.1 Source Code of *pdf.cu*

The text below is the code for the modifications made to the *pdf.cu* file from the previously named *pdf0.c*.

```
1  /*-----
2  Program pdf0.c computes a pair distribution function for n atoms
3  given the 3D coordinates of the atoms.
4  -----*/
5  #include <stdio.h>
6  #include <math.h>
7  #include <time.h>
8  #include <stdlib.h>
9  #include <cuda.h>
10
11 #define NHBIN 2000 // Histogram size
12
13 float al[3]; // Simulation box lengths
14 int n; // Number of atoms
15 float *r; // Atomic position array
16 FILE *fp;
17
18 /*
19     ADDED
20 */
21 __constant__ float DALTH[3];
22 __constant__ int DN;
23 __constant__ float DDRH;
24
25
26 // float SignR(float v, float x) {if (x > 0) return v; else return -v;}
27 /*
28     ADDED
29 */
30 __device__ float d_SignR(float v, float x)
31 {
32     if (x > 0)
33         return v;
34     else
35         return -v;
36 }
37
38 /*
39     ADDED
```

```

40  */
41  --global-- void gpu-histogram_kernel(float *r, float *nhis)
42  {
43      int i, j, a, ih;
44      float rij, dr;
45
46      int iBlockBegin = (DN / gridDim.x) * blockIdx.x;
47      int iBlockEnd = (DN / gridDim.x) * (blockIdx.x + 1);
48      if (blockIdx.x == gridDim.x - 1)
49          iBlockEnd = DN;
50      int jBlockBegin = (DN / gridDim.y) * blockIdx.y;
51      int jBlockEnd = (DN / gridDim.y) * (blockIdx.y + 1);
52      if (blockIdx.y == gridDim.y - 1)
53          jBlockEnd = DN;
54      for (i = iBlockBegin + threadIdx.x; i < iBlockEnd; i += blockDim.x)
55      {
56          for (j = jBlockBegin + threadIdx.y; j < jBlockEnd; j += blockDim.y)
57          {
58              if (i < j)
59              {
60                  // Process (i,j) atom pair
61                  rij = 0.0;
62
63                  for (a = 0; a < 3; a++)
64                  {
65                      dr = r[3 * i + a] - r[3 * j + a];
66                      /* Periodic boundary condition */
67                      dr = dr - d.SignR(DALTH[a], dr - DALTH[a]) - d.SignR(DALTH[a], dr + DALTH[a]
68                      ));
69                      rij += dr * dr;
70                  }
71                  rij = sqrt(rij); /* Pair distance */
72                  ih = rij / DDRH;
73
74                  // nhis[ih] += 1.0;
75                  atomicAdd(&nhis[ih], 1.0); /*In order to avoid race condition
76              } // end if i<j
77          } // end for j
78      } // end for i
79  }
80
81  /*-----*/
82  void histogram()
83  {
84      /*
85      PREVIOUS CODE NOT SHOWN
86      */
87
88      cudaMalloc((void**)&dev_r, sizeof(float)*3*n);
89      cudaMalloc((void**)&dev_nhis, sizeof(float)*NHBIN);
90
91      cudaMemcpy(dev_r, r, 3*n*sizeof(float), cudaMemcpyHostToDevice);
92      cudaMemset(dev_nhis, 0.0, NHBIN*sizeof(float));
93
94      cudaMemcpyToSymbol(DALTH, alth, sizeof(float)*3, 0, cudaMemcpyHostToDevice);
95      cudaMemcpyToSymbol(DN, &n, sizeof(int), 0, cudaMemcpyHostToDevice);
96      cudaMemcpyToSymbol(DDRH, &drh, sizeof(float), 0, cudaMemcpyHostToDevice);
97
98      // Compute dev_nhis on GPU: dev_r[] -> dev_nhis[]

```

```

99 dim3 numBlocks(8,8,1);
100 dim3 threads_per_block(16,16,1);
101 gpu_histogram_kernel<<<numBlocks, threads_per_block>>>(dev_r, dev_nhis);
102
103 cudaMemcpy(nhis, dev_nhis, NHBIN*sizeof(float), cudaMemcpyDeviceToHost);
104 cudaFree(dev_r);
105 cudaFree(dev_nhis);
106
107 density = n / (al[0] * al[1] * al[2]);
108 /* Print out the histogram */
109 fp = fopen("pdf.d", "w");
110 for (ih = 0; ih < NHBIN; ih++)
111 {
112     gr = nhis[ih] / (2 * M_PI * pow((ih + 0.5) * drh, 2) * drh * density * n);
113     fprintf(fp, "%e %e\n", (ih + 0.5) * drh, gr);
114 }
115 fclose(fp);
116 free(nhis);
117 }
118
119 /*-----*/
120 int main()
121 {
122     /*-----*/
123     int i;
124     float cpu1, cpu2;
125
126     /* Read the atomic position data */
127     fp = fopen("pos.d", "r");
128     fscanf(fp, "%f %f %f", &(al[0]), &(al[1]), &(al[2]));
129     fscanf(fp, "%d", &n);
130     r = (float *)malloc(sizeof(float) * 3 * n);
131     for (i = 0; i < n; i++)
132         fscanf(fp, "%f %f %f", &(r[3 * i]), &(r[3 * i + 1]), &(r[3 * i + 2]));
133     fclose(fp);
134
135     /* Compute the histogram */
136     cpu1 = ((float)clock()) / CLOCKS_PER_SEC;
137     histogram();
138     cpu2 = ((float)clock()) / CLOCKS_PER_SEC;
139     printf("Execution time (s) = %le\n", cpu2 - cpu1);
140
141     free(r);
142     return 0;
143 }

```

## 1.2 Plot of Pair Distribution Function

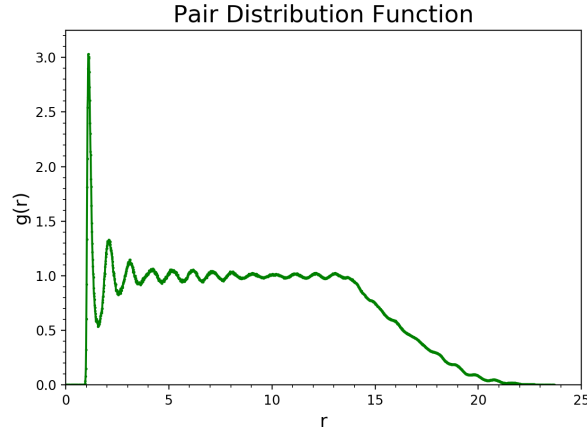


Figure 1: Plot of Pair Distribution Function

## 2 MPI+OpenMP+CUDA Computation of $\pi$

### 2.1 Source Code of *pi3.cu*

The text below is the code for the modifications made to the *pi3.cu* file from the previously named *hyypi\_setdevice.cu*.

```
1 // Hybrid MPI+OpenMP+CUDA computation of Pi
2 #include <stdio.h>
3 #include <mpi.h>
4 #include <omp.h> // NEW
5 #include <cuda.h>
6
7 #define NBIN 10000000 // Number of bins
8 #define NUM_DEVICE 2 // # of GPU devices = # of OpenMP threads // NEW
9 #define NUM_BLOCK 13 // Number of thread blocks
10 #define NUM_THREAD 192 // Number of threads per block
11
12 // Kernel that executes on the CUDA device
13 __global__ void cal_pi(float *sum, int nbin, float step, float offset, int nthreads, int
    nblocks) {
14     int i;
15     float x;
16     int idx = blockIdx.x*blockDim.x+threadIdx.x; // Sequential thread index across the
        blocks
17     for (i=idx; i<nbin; i+=nthreads*nblocks) { // Interleaved bin assignment to
        threads
18         x = offset+(i+0.5)*step;
19         sum[idx] += 4.0/(1.0+x*x);
20     }
21 }
22
23 int main(int argc, char **argv) {
```

```

24  /*
25  PREVIOUS CODE NOTE SHOWN
26  */
27
28  MPI_Init(&argc,&argv);
29  MPI_Comm_rank(MPLCOMM_WORLD,&myid); // My MPI rank
30  MPI_Comm_size(MPLCOMM_WORLD,&nproc); // Number of MPI processes
31  // nbin = NBIN/nproc; // Number of bins per MPI process
32  // step = 1.0/(float)(nbin*nproc); // Step size with redefined number of bins
33  // offset = myid*step*nbin; // Quadrature-point offset
34  //NEW
35  omp_set_num_threads(NUMDEVICE); // One OpenMP thread per GPU device
36  nbin = NBIN/(nproc*NUMDEVICE); // # of bins per OpenMP thread
37  step = 1.0/(float)(nbin*nproc*NUMDEVICE);
38
39  #pragma omp parallel private(offset, sumHost, sumDev, tid, dev_used) reduction(+:pi)
40  {
41      int mpid = omp_get_thread_num();
42      offset = (NUMDEVICE*myid+mpid)*step*nbin; // Quadrature-point offset
43      cudaSetDevice(mpid%2);
44
45      // cudaSetDevice(myid%2);
46      size_t size = NUMBLOCK*NUMTHREAD*sizeof(float); //Array memory size
47      sumHost = (float *) malloc(size); // Allocate array on host
48      cudaMalloc((void **) &sumDev, size); // Allocate array on device
49      cudaMemset(sumDev,0, size); // Reset array in device to 0
50      // Calculate on device (call CUDA kernel)
51      cal_pi <<<dimGrid,dimBlock>>> (sumDev,nbin, step, offset, NUMTHREAD, NUMBLOCK);
52      // Retrieve result from device and store it in host array
53      cudaMemcpy(sumHost, sumDev, size, cudaMemcpyDeviceToHost);
54      // Reduction over CUDA threads
55      for(tid=0; tid<NUMTHREAD*NUMBLOCK; tid++)
56          pi += sumHost[tid];
57      pi *= step;
58      // CUDA cleanup
59      free(sumHost);
60      cudaFree(sumDev);
61      cudaGetDevice(&dev_used);
62      printf("myid = %d; mpid = %d; device used = %d; partial pi = %f\n", myid, mpid,
63      dev_used, pi);
64      // printf("myid = %d; device used = %d; partial pi = %f\n",myid,dev_used,pi);
65  } // end omp parallel
66
67  // Reduction over MPI processes
68  MPI_Allreduce(&pi,&pig,1,MPL_FLOAT,MPL_SUM,MPLCOMM_WORLD);
69  if (myid==0) printf("PI = %f\n", pig);
70
71  MPI_Finalize();
72  return 0;
}

```

## 2.2 Output of *pi3.sl*

```
=====
SLURM_JOB_ID = 6305698
SLURM_JOB_NODELIST = e07-[07-08]
TMPDIR = /tmp/SLURM_6305698
=====
myid = 1; mpid = 1: device used = 1; partial pi = 0.567582
myid = 1; mpid = 0: device used = 0; partial pi = 0.719409
myid = 0; mpid = 1: device used = 1; partial pi = 0.874671
myid = 0; mpid = 0: device used = 0; partial pi = 0.979926
PI = 3.141588
```

Figure 2: Printout of pi3.sl